

# 1. Data Preprocessing & Cleaning

Before we can analyze or build models, we need to ensure our dataset is **clean, complete, and consistent**.

This step includes:

- Handling missing values
  - Removing duplicate records
  - Fixing inconsistencies and errors
  - Handling outliers
- 

## Handling Missing Values in This Dataset

### 1.1 Identifying Missing Values

Looking at the dataset, missing values are represented as "?" instead of blank spaces or NaN. This means we first need to **convert "?" into proper NaN values** so we can handle them properly.

The columns that contain missing values in this dataset are:

- "normalized-losses"
- "num-of-doors"
- "bore"
- "stroke"
- "horsepower"
- "peak-rpm"
- "price"

Now, let's discuss how we handle missing values for each column.

---

### 1.2 Suggested Methods to Handle Missing Values

Each column needs a different approach depending on **its data type and importance**.

#### 1. normalized-losses (Numeric)

- **Best approach: Median Imputation**

Since this data is numeric and could have extreme values (outliers), using the **median** is better than the mean.

- **Alternative:** Drop this column **if it's not strongly correlated** with price.

### Why not use mean imputation?

If some cars have extremely high normalized-losses, it can **pull the mean up**, making imputed values unreliable.

---

## 2. num-of-doors (Categorical)

- **Best approach: Mode Imputation** (fill missing values with the most frequent category).

Since most cars in the dataset have **four doors**, we can safely fill the missing values with "four".

- **Alternative:** Create a new category called "unknown".

---

## 3. bore & stroke (Numeric)

- **Best approach: Mean Imputation** (replace missing values with the mean).

Since these values tend to follow a normal distribution, using the mean works best.

- **Alternative:** Use a regression model to predict missing values based on similar features like engine-size or compression-ratio.

---

## 4. horsepower & peak-rpm (Numeric)

- **Best approach: Mean Imputation**

Cars of the same brand and engine size typically have similar horsepower and peak-rpm values.

- **Alternative:** Use K-Nearest Neighbors (KNN) Imputation to fill in values based on similar cars.

---

## 5. price (Numeric, Target Variable)

- **Best approach: Drop rows with missing price values**

Since price is the value we want to predict, it **cannot** be imputed.

Imputing it could introduce bias and make the model unreliable.

---

## 1.3 Alternative Methods for Handling Missing Values

If we want to try more **advanced** techniques, we can use:

- **K-Nearest Neighbors (KNN) Imputation:** Finds the most similar cars based on other features and fills in missing values.
  - **Regression Imputation:** Predicts missing values using a regression model.
- 

## Handling Duplicate Records

- **Why check for duplicates?**  
Duplicate records can skew analysis and **cause bias in predictions**.
  - **How to handle duplicates?**
    - Identify duplicate rows in the dataset.
    - If exact duplicates exist, **remove them**.
    - If near-duplicates exist (e.g., same car make & model but different pricing), **inspect them manually**.
- 

## Handling Inconsistent Data

Data inconsistencies can appear due to typos, unit mismatches, or inconsistent formatting. This dataset has some **common inconsistencies** that need fixing.

### 1.4 Fixing Data Type Issues

Several numeric columns are stored as **text (object type)**, which can cause issues in analysis. Columns affected:

- "bore"
- "stroke"
- "horsepower"
- "peak-rpm"
- "price"

**Fix:** Convert these columns to proper numerical data types.

---

### 1.5 Standardizing Categorical Variables

Certain categorical columns have **inconsistent labels** that need to be fixed.

- "num-of-cylinders" is stored as words ("four", "six") instead of numbers.
    - **Fix:** Convert to numerical representation (e.g., "four" → 4).
  - "fuel-type" might have inconsistent capitalization ("Gas" vs "gas").
    - **Fix:** Convert everything to lowercase.
- 

## Handling Outliers

Outliers are extreme values that **can distort machine learning models**.

In this dataset, outliers are likely in:

- "price"
- "horsepower"
- "engine-size"
- "curb-weight"

### How to Detect Outliers?

Method	Purpose
Boxplots	Shows extreme values visually
Z-score Method	Removes values beyond 3 standard deviations
Interquartile Range (IQR)	Removes values outside $[Q1 - 1.5 \cdot IQR, Q3 + 1.5 \cdot IQR]$

### How to Handle Outliers?

- **For price and horsepower:** Use IQR-based filtering to remove extreme values.
- **For engine-size:** Use log transformation to reduce the impact of extreme values.

## 2. Exploratory Data Analysis (EDA)

### Why is EDA Important?

Before jumping into machine learning, we need to **understand the dataset** so that we:

- Know how the data is distributed.
  - Identify correlations between variables.
  - Detect outliers and anomalies.
  - Decide which features are important.
- 

# Step 1: Univariate Analysis (Analyzing One Variable at a Time)

## 2.1 Exploring Numeric Columns

### Key Questions to Answer:

1. What is the distribution of each numeric feature?
2. Are there any outliers?
3. Are there any unusual trends (e.g., skewness)?

### How to Analyze Numeric Features?

Numeric Column	Best EDA Techniques	What It Tells Us
price	Histogram, Boxplot, KDE Plot	Distribution, skewness, outliers
horsepower	Histogram, Boxplot, Violin Plot	Whether certain horsepower values are common/uncommon
engine-size	Histogram, KDE Plot	Which engine sizes are most common
curb-weight	Histogram, Boxplot	If weight is normally distributed or skewed

compression -ratio	Histogram, KDE Plot	If some cars have extreme compression ratios
-----------------------	---------------------	--

Key Insights We Expect to Find

- price is likely **right-skewed** (more cheaper cars, fewer expensive ones).
  - horsepower might show **bimodal distribution** (different peaks for economy vs sports cars).
  - engine-size may be **normally distributed**, but with some extreme values.
- 

2.2 Exploring Categorical Columns

Key Questions to Answer:

1. Which categories are most/least common?
2. Are there any rare categories?
3. Do we need to group similar categories together?

How to Analyze Categorical Features?

Categorical Column	Best EDA Techniques	What It Tells Us
fuel-type	Bar Chart, Pie Chart	Are most cars gasoline or diesel?
num-of-doors	Bar Chart, Count Plot	Are most cars 2-door or 4-door?
body-style	Bar Chart, Pie Chart	What car body styles are most popular?
drive-wheels	Bar Chart, Count Plot	How common are FWD, RWD, and 4WD cars?

engine-type	Bar Chart, Count Plot	Which engine types dominate?
-------------	-----------------------	------------------------------

Key Insights We Expect to Find

- Most cars in the dataset are likely **gasoline-powered**.
  - num-of-doors might have more **four-door** cars.
  - drive-wheels might show **more FWD cars** since most economy cars are FWD.
- 

# Step 2: Bivariate Analysis (Analyzing Relationships Between Two Variables)

Now, we explore **how two variables interact**.

## 2.3 Numeric vs. Numeric Relationships

Key Questions to Answer:

1. How do numerical features relate to price?
2. Are there strong correlations?
3. Are relationships linear or nonlinear?

### How to Analyze Numeric vs. Numeric Relationships?

Pair of Variables	Best EDA Techniques	What It Tells Us
engine-size vs price	Scatter Plot, Regression Plot	Larger engines should increase price
horsepower vs price	Scatter Plot, Correlation Matrix	More horsepower = Higher price?

curb-weight vs price	<b>Scatter Plot, Correlation Matrix</b>	Heavier cars may cost more
city-mpg vs highway-mpg	<b>Scatter Plot, Regression Plot</b>	Directly correlated (higher city mpg → higher highway mpg)

### Expected Insights

- engine-size and horsepower should be **strongly correlated** with price.
  - curb-weight may also have a **positive correlation** with price.
  - city-mpg and highway-mpg should be **strongly correlated**, meaning that cars with good city mileage tend to have good highway mileage too.
- 

## 2.4 Categorical vs. Numeric Relationships

### Key Questions to Answer:

1. Do categorical features affect price?
2. Do certain categories result in higher/lower prices?

### How to Analyze Categorical vs. Numeric Relationships?

Categorical Variable	Numeric Variable	Best EDA Techniques	What It Tells Us
fuel-type	price	<b>Boxplot, Violin Plot</b>	Do diesel cars cost more?
body-style	price	<b>Boxplot, Bar Chart</b>	Which car body types are expensive?
num-of-doors	price	<b>Boxplot</b>	Are 4-door cars more expensive?



drive-wheels	price	Boxplot, Violin Plot	Are AWD/4WD cars more expensive?
--------------	-------	----------------------	----------------------------------

### Expected Insights

- fuel-type: Diesel cars **might** be more expensive than gasoline.
  - body-style: Convertibles or luxury sedans might have **higher prices**.
  - drive-wheels: AWD/4WD vehicles might be **pricier than FWD cars**.
- 

## 2.5 Categorical vs. Categorical Relationships

### Key Questions to Answer:

1. Are some categories related to each other?
2. Are some categories more common together?

### How to Analyze Categorical vs. Categorical Relationships?

Categorical Variable 1	Categorical Variable 2	Best EDA Techniques	What It Tells Us
body-style	drive-wheels	Stacked Bar Chart, Heatmap	Do sedans mostly have FWD?
fuel-type	engine-type	Heatmap, Crosstab	Do certain engines run on specific fuel types?

### Expected Insights

- Most **sedans** are likely to be **FWD**.
  - Most **sports cars** are likely **RWD**.
-

# Step 3: Multivariate Analysis (Analyzing More Than Two Variables)

Now, we explore how **multiple features interact**.

## 2.6 Correlation Analysis (Multiple Numeric Features)

- A **Correlation Matrix (Heatmap)** helps us see which numeric features are strongly related.
- **Expected High Correlations**
  1. engine-size vs. horsepower
  2. horsepower vs. price
  3. curb-weight vs. engine-size

## 2.7 Dimensionality Reduction (PCA)

- If we have **too many correlated features**, we can use **Principal Component Analysis (PCA)** to reduce dimensions while keeping most of the information.
- 

## Key Takeaways from EDA

1. price is likely **right-skewed**, meaning there are many cheaper cars and fewer expensive ones.
2. engine-size and horsepower should be **strongly correlated** with price.
3. drive-wheels and fuel-type may have a **big impact on price**.
4. city-mpg and highway-mpg are likely **strongly correlated**.
5. body-style and num-of-doors may have **no strong relationship** with price.

## 3. Machine Learning Algorithm Selection

Now that we've explored and understood the dataset through **EDA**, we can move on to **choosing the right machine learning models** for different possible tasks.

Since this dataset contains information about **cars and their prices**, there are multiple **machine learning problems** we can solve:

1. **Regression** (Predicting a continuous value like price)
2. **Classification** (Predicting a category like fuel-type)
3. **Unsupervised Learning** (Clustering similar cars together)

---

## Step 1: Define the Machine Learning Tasks

Before selecting models, we need to decide **what we want to predict**.

### 3.1 Regression Task: Predicting price

- **Goal:** Given car features (e.g., engine-size, horsepower, drive-wheels), predict the price of a car.
- **Why Regression?**
  - price is a **continuous numerical value**, making regression the right choice.

### 3.2 Classification Task: Predicting fuel-type

- **Goal:** Predict whether a car runs on "gas" or "diesel" based on its features.
- **Why Classification?**
  - fuel-type is a **categorical variable** (gas or diesel), making classification appropriate.

### 3.3 Clustering Task: Grouping Similar Cars

- **Goal:** Automatically group similar cars together based on engine-size, horsepower, and curb-weight.
- **Why Unsupervised Learning?**
  - This task doesn't have predefined labels, making **clustering** the best approach.

---

## Step 2: Selecting the Right Machine Learning Models

Now that we know what tasks to solve, let's choose the best algorithms.

---

### 3.1 Best Algorithms for Regression (Predicting price)

Since price is a **continuous variable**, we use **regression models**.

Algorithm	When to Use?	Strengths	Weaknesses
<b>Linear Regression</b>	If relationships are linear	Simple, interpretable	Struggles with nonlinearity
<b>Polynomial Regression</b>	If price depends on higher-order relationships	Captures nonlinear patterns	Can overfit on small datasets
<b>Decision Tree Regression</b>	If price has many complex relationships	Handles nonlinear relationships	Can overfit if not pruned
<b>Random Forest Regression</b>	When high accuracy is needed	Reduces overfitting by using multiple trees	Slower than simple models
<b>Gradient Boosting (XGBoost, LightGBM)</b>	If maximum performance is needed	Handles large datasets well, best accuracy	Computationally expensive

Which model is best?

- **Start with Linear Regression** to check for simple relationships.
- If relationships **aren't linear**, try **Random Forest**.
- For the **best performance**, use **XGBoost or LightGBM**.

### 3.2 Best Algorithms for Classification (Predicting fuel-type)

Since fuel-type is a **categorical variable (gas or diesel)**, we use **classification models**.

Algorithm	When to Use?	Strengths	Weaknesses
-----------	--------------	-----------	------------

<b>Logistic Regression</b>	When the relationship is mostly linear	Simple, interpretable	Poor with complex data
<b>Decision Tree Classifier</b>	When handling categorical data	Easy to interpret	Can overfit
<b>Random Forest Classifier</b>	If accuracy is more important	Reduces overfitting	Slower for large datasets
<b>Support Vector Machine (SVM)</b>	If fuel-type is hard to separate	Works well for small datasets	Computationally expensive
<b>Gradient Boosting (XGBoost, LightGBM)</b>	If the best accuracy is needed	Best for complex patterns	Requires tuning

Which model is best?

- Start with **Logistic Regression** for a simple, explainable model.
- Try **Random Forest** if there are **nonlinear relationships**.
- If **performance is the priority**, use **XGBoost**.

### 3.3 Best Algorithms for Clustering (Grouping Similar Cars)

If we want to automatically group similar cars together, we use **unsupervised learning**.

Algorithm	When to Use?	Strengths	Weaknesses
<b>K-Means Clustering</b>	If we assume clusters have similar sizes	Simple and fast	Requires choosing k (number of clusters)

<b>Hierarchical Clustering</b>	If we want a hierarchy of clusters	No need to choose k	Computationally expensive
<b>DBSCAN</b>	If we expect different cluster sizes	Handles noise well	Can struggle with high-dimensional data

Which model is best?

- **Start with K-Means**, setting  $k = 3$  (e.g., economy, mid-range, luxury cars).
- If clusters aren't well-separated, try **Hierarchical Clustering**.

## Step 3: Evaluating Model Performance

Once we train our models, we must **evaluate their accuracy** using appropriate metrics.

### 3.4 Best Metrics for Regression (Predicting price)

Metric	What It Measures	When to Use
<b>Mean Absolute Error (MAE)</b>	Average error in dollars	If we care about <b>actual dollar errors</b>
<b>Mean Squared Error (MSE)</b>	Penalizes large errors	If large errors should be minimized
<b>Root Mean Squared Error (RMSE)</b>	Similar to MSE but in original units	If we want errors in the same scale as price
<b>R<sup>2</sup> Score</b>	How well the model explains variance	To see <b>overall model fit</b>

- ◆ **Best Choice for This Dataset:** Use **RMSE** (since price is in dollars) and **R<sup>2</sup> Score** (to measure overall accuracy).
- 

### 3.5 Best Metrics for Classification (Predicting fuel-type)

Metric	What It Measures	When to Use
Accuracy	Overall correctness	When classes are balanced
Precision	How many predicted "diesel" cars are actually diesel	When false positives are costly
Recall	How many actual diesel cars were correctly predicted	When missing a "diesel" car is costly
F1-Score	Balance between precision & recall	When data is imbalanced

- ◆ **Best Choice for This Dataset:** Use **Accuracy** first, then check **F1-score** if the classes are imbalanced.
- 

### 3.6 Best Metrics for Clustering

Metric	What It Measures	When to Use
Silhouette Score	How well clusters are separated	General clustering evaluation
Davies-Bouldin Index	Measures compactness of clusters	When clusters are overlapping

- ◆ **Best Choice for This Dataset:** Use **Silhouette Score** to evaluate cluster quality.
- 

## Final Summary of Machine Learning Model Selection

Task	Best Models	Best Evaluation Metric
Regression (price)	Random Forest, XGBoost	RMSE, R <sup>2</sup> Score
Classification (fuel-type)	Logistic Regression, Random Forest	Accuracy, F1-Score
Clustering (grouping cars)	K-Means, Hierarchical	Silhouette Score

## 4. Feature Engineering & Model Optimization

Now that we have selected the best machine learning models for each task, we need to **improve model performance** through **feature engineering** and **hyperparameter tuning**.

This step involves:

1. **Transforming features to improve model accuracy.**
  2. **Selecting the most important features to avoid overfitting.**
  3. **Tuning hyperparameters to optimize model performance.**
  4. **Using advanced techniques like ensemble learning for better results.**
- 

### Step 1: Feature Engineering

Feature engineering involves **modifying, creating, or selecting** features to help machine learning models perform better.

#### 4.1 Feature Engineering for Numeric Features



Numeric features may need **scaling, transformations, or new feature creation**.

### Handling Skewed Numeric Features

- Some numeric features in this dataset (e.g., price, horsepower) are **right-skewed**.
- We can **apply transformations** to make them **more normally distributed**.

Feature	Problem	Best Transformation
price	Right-skewed	<b>Log Transformation</b>
horsepower	Right-skewed	<b>Log Transformation</b>
compression-ratio	Outliers present	<b>Standardization (Z-score scaling)</b>

#### ♦ Why use log transformation?

- It **reduces the impact of outliers** and makes relationships **more linear** for models like **Linear Regression**.
  - For example, instead of working with **actual price values**, we use **log(price)**.
- 

### Scaling Numeric Features

- Models like **Linear Regression, SVM, and KNN** perform **better** when numeric values are scaled.
- Scaling ensures that **all numeric values are on the same scale**, preventing features with large values (e.g., curb-weight) from dominating the model.

Scaling Method	When to Use?	Best Features to Apply
<b>Min-Max Scaling</b>	When we want values between 0 and 1	engine-size, horsepower, curb-weight
<b>Standardization (Z-score Scaling)</b>	When data is normally distributed	compression-ratio, city-mpg, highway-mpg

#### ♦ Which one to use?

- **Min-Max Scaling** is best for **Tree-based models** (Random Forest, XGBoost).
- **Z-score Scaling** is best for **Linear models (SVM, Logistic Regression, Linear Regression)**.

---

## 4.2 Feature Engineering for Categorical Features

Since machine learning models **do not understand text**, we need to convert categorical variables into **numeric format**.

Categorical Feature	Encoding Type	When to Use?
fuel-type	One-Hot Encoding	When categories are <b>not ordinal</b>
drive-wheels	One-Hot Encoding	When categories are <b>not ordinal</b>
num-of-cylinders	Label Encoding	If the number of cylinders is <b>ordinal</b>
body-style	One-Hot Encoding	When categories <b>don't have order</b>

### ♦ Why use One-Hot Encoding?

- It **creates separate binary columns** (e.g., fuel-type\_gas, fuel-type\_diesel).
- Works best for models like **Linear Regression** and **Neural Networks**.

### ♦ Why use Label Encoding?

- It **converts categories into numbers** (e.g., four cylinders → 4).
- Works best when the categories have a **natural order**.

---

## 4.3 Creating New Features

Sometimes, **new features** can improve model accuracy.

New Feature	How to Calculate	Why it Helps?
power-to-weight-ratio	horsepower / curb-weight	Helps predict performance
mpg-difference	highway-mpg - city-mpg	Shows efficiency difference

luxury-score	engine-size * horsepower	Creates a score for expensive cars
--------------	--------------------------	------------------------------------

- ◆ **Why create new features?**
    - Machine learning models **perform better** when given features that better represent real-world patterns.
- 

#### 4.4 Feature Selection (Keeping Only the Best Features)

We should **remove unnecessary features** to:

- **Reduce overfitting** (less noise in the model).
- **Improve speed** (models train faster).
- **Increase accuracy** (no irrelevant data).

Feature Selection Method	How It Works	Best For
Correlation Analysis	Removes highly correlated features	Regression
Recursive Feature Elimination (RFE)	Keeps the most important features	Tree-based models
Mutual Information	Measures how much each feature affects the target variable	Classification

- ◆ **Example: Removing Highly Correlated Features**
    - If engine-size and curb-weight have a **correlation > 0.9**, we **drop one** to avoid redundancy.
- 

### Step 2: Model Optimization (Hyperparameter Tuning)

Once we have **engineered the best features**, we need to **fine-tune the models** to improve performance.

#### 4.5 Why Hyperparameter Tuning?

- Each machine learning model has **settings (hyperparameters)** that control how it learns.

- If not optimized, models may **overfit or underperform**.
- 

## 4.6 Hyperparameter Tuning Techniques

Tuning Method	How It Works	Best Use Case
<b>Grid Search (GridSearchCV)</b>	Tries all possible hyperparameter combinations	Best when we have <b>few parameters</b>
<b>Random Search (RandomizedSearchCV)</b>	Randomly selects hyperparameters from a range	Best for <b>fast tuning</b>
<b>Bayesian Optimization</b>	Learns which hyperparameters are best over time	Best for <b>deep learning &amp; complex models</b>

### ♦ Example: Tuning a Random Forest Model

- Instead of using **default settings**, we fine-tune:
    - `n_estimators` (number of trees)
    - `max_depth` (depth of trees)
    - `min_samples_split` (minimum samples per split)
- 

## 4.7 Ensemble Learning (Combining Multiple Models)

- Sometimes, instead of using **one model**, we combine **multiple models** to **increase accuracy**.

Ensemble Method	How It Works	Best For
<b>Bagging (Random Forest)</b>	Uses multiple weak models (trees) to reduce overfitting	Large datasets
<b>Boosting (XGBoost, LightGBM)</b>	Corrects mistakes of previous models	Complex relationships
<b>Stacking</b>	Combines predictions from multiple models	Best accuracy

### ♦ Which one should we use?

- **Bagging** is good when data **varies a lot** (Random Forest).
- **Boosting** works best when we need **high accuracy** (XGBoost).

- **Stacking** is **powerful** but requires multiple models.

---

## Final Summary of Feature Engineering & Model Optimization

Task	Best Technique	Why?
Scaling Numeric Features	Min-Max Scaling (Tree Models), Standardization (Linear Models)	Prevents large values from dominating
Encoding Categorical Features	One-Hot Encoding, Label Encoding	Converts text into numbers
Feature Selection	Correlation Analysis, RFE	Removes irrelevant data
Hyperparameter Tuning	Grid Search, Random Search	Improves model accuracy
Ensemble Learning	Random Forest, XGBoost	Increases prediction power

---

## 5. Model Deployment & Real-World Considerations

Now that we have built and optimized the best models, it's time to **deploy them** and consider how they will perform in the real world.

This step ensures that our model can:

- **Be accessed by users** through a web app, API, or cloud service.
- **Handle real-world data efficiently** (including unseen data).
- **Stay accurate over time** through monitoring and retraining.

---

### Step 1: Choosing a Deployment Method

There are multiple ways to deploy a machine learning model depending on the **use case** and **requirements**.

## 5.1 Best Deployment Methods for This Dataset

Deployment Method	How It Works	Best Use Case
Web API (Flask/FastAPI)	Turns the model into an API that can be accessed online	If the model needs to be used by multiple applications
Cloud Deployment (AWS, GCP, Azure)	Deploys the model on cloud servers for scalability	If many users need to access the model
Edge Deployment (IoT/Embedded Systems)	Runs the model on small devices	If used in a car dashboard or IoT system
Streamlit Web App	Provides a simple web interface for predictions	If an interactive UI is needed for users

### ♦ Best choice for this dataset?

- If we want users to **send car features and get a price prediction**, a **Flask/FastAPI API** is ideal.
  - If we need a **user-friendly UI**, a **Streamlit web app** can display results in real time.
  - If we expect **thousands of users**, cloud deployment (AWS, GCP) is better.
- 

## Step 2: Handling Real-World Data

Once deployed, the model will **receive real-world data**, which may be **messy or different** from training data.

## 5.2 Common Real-World Data Challenges

Challenge	How to Handle It
Missing values in new data	Use the same imputation methods as training
Categorical values not seen during training	Use a "Unknown" category for new values
Outliers in new data	Apply the same outlier detection methods as training
Changing data distributions (Concept Drift)	Continuously monitor and retrain the model

◆ **Why is this important?**

- If a **new car brand appears**, the model should still work without failing.
  - If **fuel types change** (e.g., electric cars become common), the model needs updating.
- 

**Step 3: Model Monitoring & Retraining**

Over time, the model’s accuracy will **decline** due to **changes in real-world data**.  
For example:

- New car models might have **different price patterns**.
- Fuel types might **change over time**.

**5.3 How to Monitor Model Performance?**

Monitoring Technique	What It Checks
Track Prediction Accuracy	Regularly check RMSE or F1-score on new data
Drift Detection	Check if the new data distribution is different from training data
User Feedback	Ask users if predictions seem reasonable

◆ **How often should we retrain the model?**

- If the model is used **daily**, check for drift **weekly** and retrain **monthly**.
  - If used for **historical car price estimation**, retraining **yearly** may be enough.
- 

**Step 4: Handling Data Drift**

Data drift happens when **the data distribution changes over time**, making the model outdated.

**5.4 Types of Data Drift**

Type	What Happens?	Example in This Dataset
------	---------------	-------------------------

<b>Concept Drift</b>	The relationship between features and target changes	The same engine size no longer affects price the same way
<b>Feature Drift</b>	The distribution of a feature changes	More electric cars appear, making fuel-type less relevant
<b>Label Drift</b>	The target variable distribution changes	Car prices increase due to inflation

### 5.5 Solutions to Data Drift

Solution	How It Works
<b>Scheduled Retraining</b>	Update the model every X months with new data
<b>Online Learning</b>	Continuously adjust the model with new data
<b>Adaptive Models</b>	Use models that learn from real-time feedback

- ♦ Best solution for this dataset?
- If prices change frequently, retrain the model **every 6 months**.
  - If new car features emerge, track **feature drift** and update encoding methods.

---

## Step 5: Explainability & Ethical Considerations

In real-world applications, **users need to trust the model’s predictions**.

### 5.6 Making Predictions Explainable

Explainability Method	How It Helps
<b>Feature Importance (SHAP values)</b>	Shows which factors impact predictions the most
<b>Partial Dependence Plots</b>	Shows how price changes with each feature
<b>LIME (Local Interpretable Model-Agnostic Explanations)</b>	Explains individual predictions



♦ **Example in this dataset:**

- If the model **predicts a car's price as \$30,000**, SHAP values can show **why** (e.g., engine-size contributed most).
- 

## 5.7 Ethical Considerations

Machine learning models **must be fair and unbiased**.

Ethical Concern	Solution
Bias against certain brands	Ensure make does not heavily influence price
Transparency	Show why a price is predicted
Data Privacy	Do not store user data unnecessarily

♦ **How to ensure fairness?**

- **Test the model** on different car brands to check for bias.
  - **Ensure predictions** are explainable to users.
-